

Introduction to the RISC-V ISA

Rishiyur S. Nikhil
v2022-12-08.1



These slides (with sources) are available at: https://github.com/rsnikhil/RISC-V_Intro

Disclaimer: any errors/opinions herein should be attributed only to the author and not to RISC-V International or to Bluespec, Inc.

Contents

- RISC-V: general information
- RISC-V ISA's modular organization
 - I(nteger) + optional extensions M, A, F, D, C, ...
 - Unprivileged and Privileged
- Unprivileged ISA overview (I, M, A, F, D, C)
- Memory and I/O
- Privileged ISA
 - CSRs (Control and Status registers)
 - Exceptions (traps and interrupts)
 - Virtual Memory in Supervisor and User privilege levels
- Common artefacts attached to cores:
 - Software Interruptor, Timer with timer interrupts
 - PLIC: Platform-Level Interrupt controller
 - Debug Module
- Comments on RISC-V software ecosystem
- Verification
 - Formal Specification of RISC-V ISA
 - Golden reference models
 - Standard ISA tests
 - Other tests

The term *RISC-V*, per se, is “only” an ISA (Instruction Set Architecture)

- Only a specification, a document, describing:
 - “Architecturally-visible” state: registers (PC, integer, PC, floating point, CSRs)
 - Repertoire of instructions, and binary representation of each one
 - Semantics (meaning) of each instruction: how it accesses/changes architecturally visible state*i.e.*, it’s the abstract view of the machine targeted by general-purpose compilers, or the Assembly Languages programmer’s view
- Not an actual CPU, processor, chip, core, ...
- *Not within the purview of the ISA*: micro-architectural choices in particular implementations, such as pipeline structure, instruction pre-fetch, branch-prediction and other speculation, bypassing, superscalarity, out-of-order execution, multithreading, caches, MHz, CPI, energy consumption, *etc.*

There can be a variety of implementations varying widely on these dimensions, from IoT microcontrollers to server-class cores to supercomputing cores.

This Introduction is only about the ISA, not any specific RISC-V implementation

Primary References

Unprivileged and Privileged ISAs:

- *The RISC-V Instruction Set Manual; Volume I: Unprivileged ISA, Document Version 20191213*, Andrew Waterman and Krste Asanović (editors), December 13, 2019, 238 pp.
- *The RISC-V Instruction Set Manual; Volume II: Privileged Architecture, Document Version 20211203*, Andrew Waterman, Krste Asanović¹ and John Hauser (editors) December 4, 2021, 155 pp.

PDFs for all RISC-V specs can be found at:

<https://riscv.org/technical/specifications/>

Clean-slate design

RISC-V was created circa. 2010, at University of California, Berkeley, USA, by a team with 4-5 decades of experience in ISA design and implementations (with colleagues who are experts in operating systems, programming languages and compilers, and software tools and ecosystems).

- U.C.Berkeley created four previous RISC designs starting in the 1980s (hence the “V” in “RISC-V”), including the parent design of the **SPARC** ISA (commercialized by Sun Microsystems, Oracle, Fujitsu, ...).

Result:

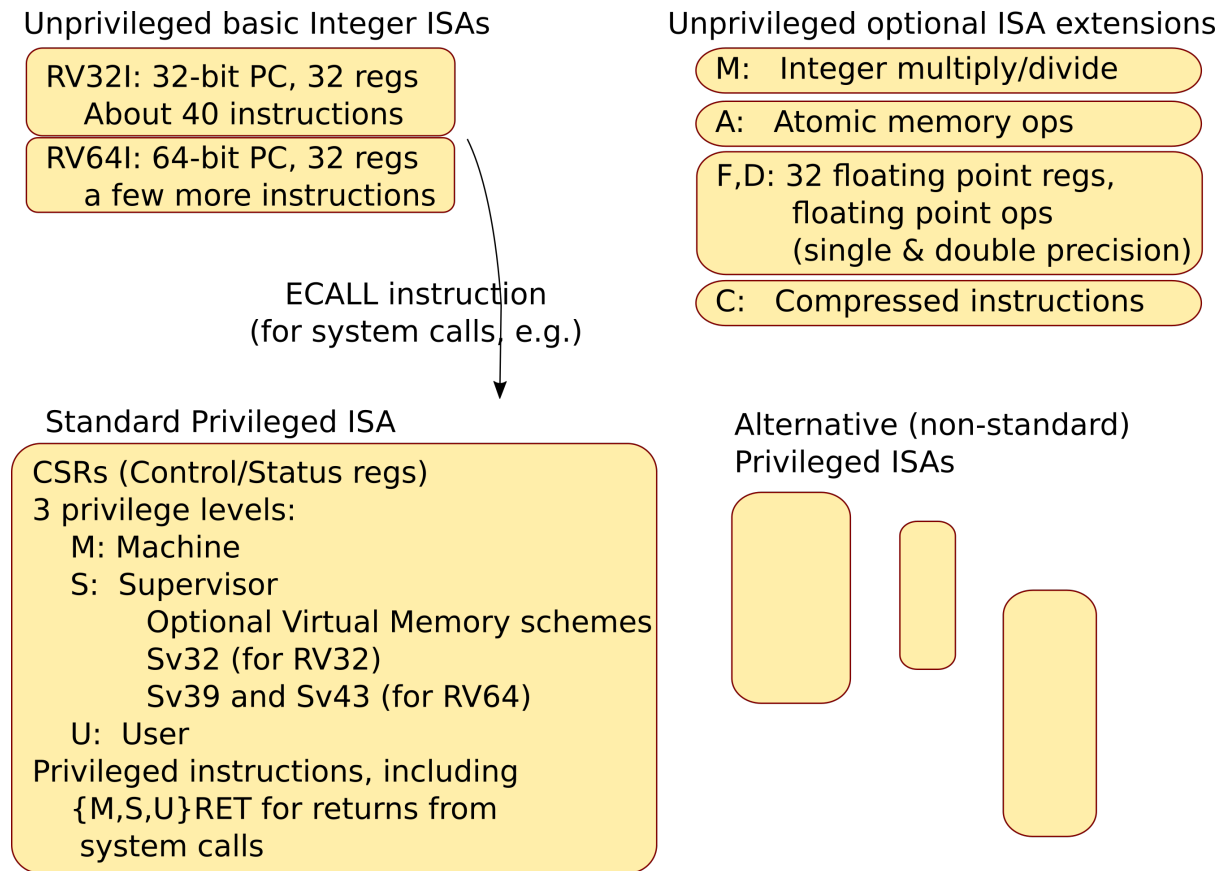
- *Very* clean, orthogonal, simple design.
- No legacy bloat (*e.g.*, branch-delay slots, ...)
- Compare: ARM’s 64-bit architecture (ARM v8-A) spec is over 6000 pages long.

RISC-V is an *Open* ISA, unlike other ISAs

- There are many well-known, mature ISAs—x86, ARMv8, Power, Sparc, ...—but they are all *proprietary*, requiring licenses and fees if you want to implement them.
- RISC-V does not require any license fee to implement.¹
 - This is central to all the current commercial interest in RISC-V.
 - This is central to *architecture researchers* who wish to innovate on state-of-the-art implementations, and indeed was the original motivation behind the creation of RISC-V by a team at Univ. of California, Berkeley.
- Because it is *open*, and already relatively mature, RISC-V seems well on its way to becoming one of the *three dominant ISAs* (along with x86 and ARM). Over time, it may also displace the numerous lesser-known ISAs used for microcontrollers.

¹ The name “RISC-V” is owned by RISC-V International, and you need their permission to use the name in commercial naming of products.

The RISC-V ISA has a *Modular* Organization



PDFs for all RISC-V specs can be found at: <https://riscv.org/technical/specifications/>

- Many other standard extensions exist: vector, crypto, bit-manipulation, ...
- Implementors can choose according to target, from small IoT/Embedded (*e.g.*, RV32IC “bare-metal”) to server-class/HPC (*e.g.*, RV64IMAFDC with vector, crypto, bit-manipulation and standard Privileged ISA).
- HW facilities allow SW to discover the configuration on which it is running.

Instruction formats

PDFs for RISC-V specs can be found at: <https://riscv.org/technical/specifications/>

- All instructions are 32-bits wide, for both RV32 and RV64.
- There are very few instruction formats (simplifies hardware-decoder in CPU pipelines):

130

Volume I: RISC-V Unprivileged ISA V20191213

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
funct7				rs2		rs1		funct3		rd		opcode		R-type
imm[11:0]						rs1		funct3		rd		opcode		I-type
imm[11:5]				rs2		rs1		funct3		imm[4:0]		opcode		S-type
imm[12 10:5]				rs2		rs1		funct3		imm[4:1 11]		opcode		B-type
imm[31:12]										rd		opcode		U-type
imm[20 10:1 11 19:12]										rd		opcode		J-type

- In the optional “C” extension (“compressed”), instructions are 16-bits wide, for small-footprint systems (IoT/Embedded/Edge).
 - Each 16-bit C instruction expands to a standard 32-bit RV32/64 instruction; so only needs hardware at the front-end of the CPU pipeline (16-bit fetch and expansion).

Base Unprivileged Integer ISA (with 32 32-bit or 64-bit registers)

PDFs for RISC-V specs can be found at: <https://riscv.org/technical/specifications/>

130

Volume I: RISC-V Unprivileged ISA V20191213

31	27	26	25	24	20	19	15	14	12	11	7	6	0	
funct7				rs2		rs1		funct3		rd		opcode		R-type
imm[11:0]				rs1		funct3		rd		opcode				I-type
imm[11:5]				rs2		rs1		funct3		imm[4:0]		opcode		S-type
imm[12 10:5]				rs2		rs1		funct3		imm[4:1 11]		opcode		B-type
				imm[31:12]						rd		opcode		U-type
				imm[20 10:1 11 19:12]						rd		opcode		J-type

RV32I Base Instruction Set

imm[31:12]				rd		0110111		LUI				
imm[31:12]				rd		0010111		AUIPC				
imm[20 10:1 11 19:12]				rd		1101111		JAL				
imm[11:0]				rs1		000		rd	1100111	JALR		
imm[12 10:5]		rs2		rs1		000		imm[4:1 11]	1100011	BEQ		
imm[12 10:5]		rs2		rs1		001		imm[4:1 11]	1100011	BNE		
imm[12 10:5]		rs2		rs1		100		imm[4:1 11]	1100011	BLT		
imm[12 10:5]		rs2		rs1		101		imm[4:1 11]	1100011	BGE		
imm[12 10:5]		rs2		rs1		110		imm[4:1 11]	1100011	BLTU		
imm[12 10:5]		rs2		rs1		111		imm[4:1 11]	1100011	BGEU		
imm[11:0]				rs1		000		rd	0000011	LB		
imm[11:0]				rs1		001		rd	0000011	LH		
imm[11:0]				rs1		010		rd	0000011	LW		
imm[11:0]				rs1		100		rd	0000011	LBU		
imm[11:0]				rs1		101		rd	0000011	LHU		
imm[11:5]		rs2		rs1		000		imm[4:0]	0100011	SB		
imm[11:5]		rs2		rs1		001		imm[4:0]	0100011	SH		
imm[11:5]		rs2		rs1		010		imm[4:0]	0100011	SW		
imm[11:0]				rs1		000		rd	0010011	ADDI		
imm[11:0]				rs1		010		rd	0010011	SLTI		
imm[11:0]				rs1		011		rd	0010011	SLTIU		
imm[11:0]				rs1		100		rd	0010011	XORI		
imm[11:0]				rs1		110		rd	0010011	ORI		
imm[11:0]				rs1		111		rd	0010011	ANDI		
0000000		shamt		rs1		001		rd	0010011	SLLI		
0000000		shamt		rs1		101		rd	0010011	SRLI		
0100000		shamt		rs1		101		rd	0010011	SRAI		
0000000		rs2		rs1		000		rd	0110011	ADD		
0100000		rs2		rs1		000		rd	0110011	SUB		
0000000		rs2		rs1		001		rd	0110011	SLL		
0000000		rs2		rs1		010		rd	0110011	SLT		
0000000		rs2		rs1		011		rd	0110011	SLTU		
0000000		rs2		rs1		100		rd	0110011	XOR		
0000000		rs2		rs1		101		rd	0110011	SRL		
0100000		rs2		rs1		101		rd	0110011	SRA		
0000000		rs2		rs1		110		rd	0110011	OR		
0000000		rs2		rs1		111		rd	0110011	AND		
fm		pred		succ		rs1		000		rd	0001111	FENCE
000000000000				00000		000		00000		1110011		ECALL
000000000001				00000		000		00000		1110011		EBREAK

- RV32I ISA (Unprivileged, 32 x 32-bit registers, integer only) has a mere 40 instructions (left).
- RV64I ISA (32 x 64-bit registers) adds 15 more instructions (below).

RV64I Base Instruction Set (in addition to RV32I)							
imm[11:0]		rs1	110	rd	0000011	LWU	
imm[11:0]		rs1	011	rd	0000011	LD	
imm[11:5]		rs2	rs1	011	imm[4:0]	0100011	SD
000000		shamt	rs1	001	rd	0010011	SLLI
000000		shamt	rs1	101	rd	0010011	SRLI
010000		shamt	rs1	101	rd	0010011	SRAI
imm[11:0]		rs1	000	rd	0011011	ADDIW	
0000000		shamt	rs1	001	rd	0011011	SLLIW
0000000		shamt	rs1	101	rd	0011011	SRLIW
0100000		shamt	rs1	101	rd	0011011	SRAIW
0000000		rs2	rs1	000	rd	0111011	ADDW
0100000		rs2	rs1	000	rd	0111011	SUBW
0000000		rs2	rs1	001	rd	0111011	SLLW
0000000		rs2	rs1	101	rd	0111011	SRLW
0100000		rs2	rs1	101	rd	0111011	SRAW

The “W” instructions operate on 32-bit data in 64-bit registers.

- AUIPC (Add Upper Immediate to PC): enables position-independent code.
- Pure “Load-Store” ISA, *i.e.*, complete separation of memory access instructions from all other instructions (Lx, Sx).
- All I/O is via memory-mapped registers; no separate instructions.
- ECALL: (Environment Call) “call-out” to Privileged ISA. Details are part of Privileged Spec.
- EBREAK: “call-out” to debugging environment (unspecified).

Possibly the cleanest, most orthogonal, industrial-strength ISA ever.

“M” extension: Integer Multiply and Divide instructions

RV32M Standard Extension						
0000001	rs2	rs1	000	rd	0110011	MUL
0000001	rs2	rs1	001	rd	0110011	MULH
0000001	rs2	rs1	010	rd	0110011	MULHSU
0000001	rs2	rs1	011	rd	0110011	MULHU
0000001	rs2	rs1	100	rd	0110011	DIV
0000001	rs2	rs1	101	rd	0110011	DIVU
0000001	rs2	rs1	110	rd	0110011	REM
0000001	rs2	rs1	111	rd	0110011	REMU
RV64M Standard Extension (in addition to RV32M)						
0000001	rs2	rs1	000	rd	0111011	MULW
0000001	rs2	rs1	100	rd	0111011	DIVW
0000001	rs2	rs1	101	rd	0111011	DIVUW
0000001	rs2	rs1	110	rd	0111011	REMW
0000001	rs2	rs1	111	rd	0111011	REMUW

- S: signed; U: Unsigned: enables signed and unsigned multiplications.
- H: upper half of double-width data: enables 64-bit multiplication-results on RV32, 128-bit multiplication-results on RV64.
- The “W” instructions operate on 32-bit data in 64-bit registers.

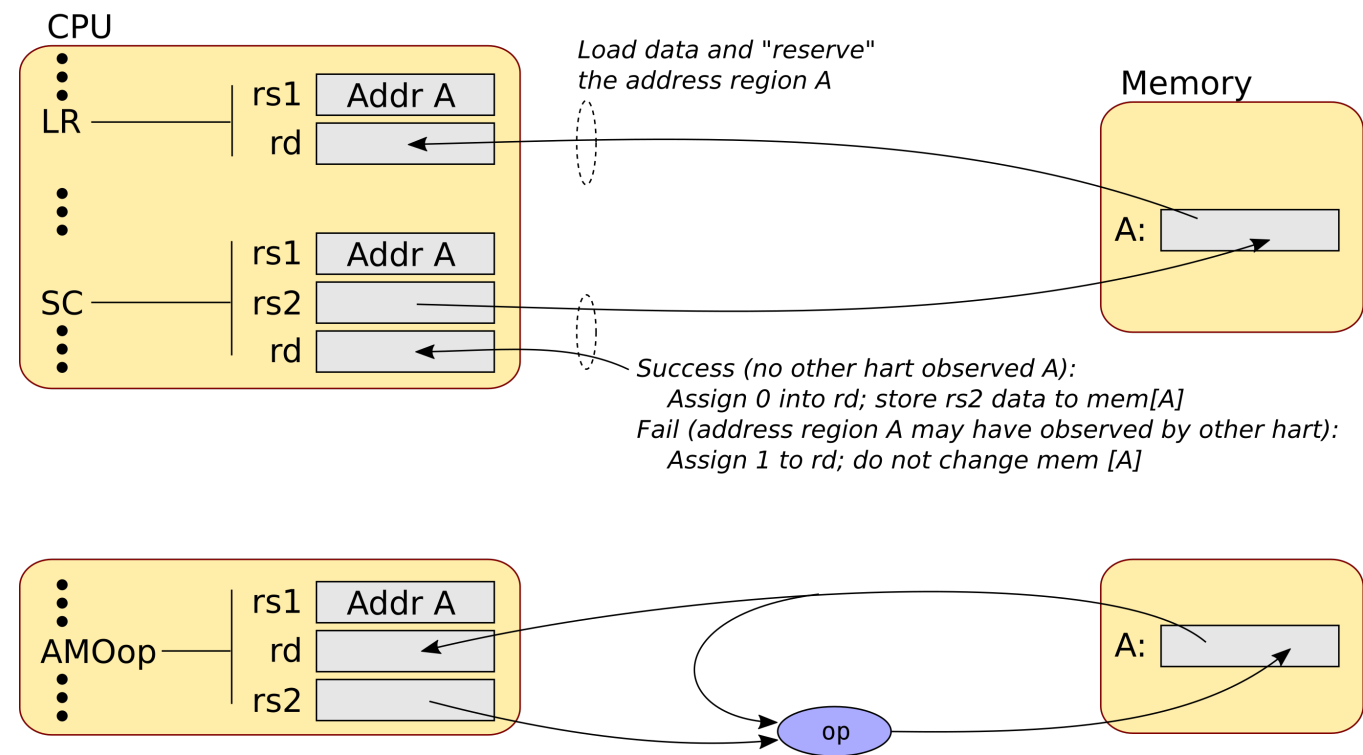
“A” extension: Atomic memory operations

RV32A Standard Extension								
00010	aq	rl	00000	rs1	010	rd	0101111	LR.W
00011	aq	rl	rs2	rs1	010	rd	0101111	SC.W
00001	aq	rl	rs2	rs1	010	rd	0101111	AMOSWAP.W
00000	aq	rl	rs2	rs1	010	rd	0101111	AMOADD.W
00100	aq	rl	rs2	rs1	010	rd	0101111	AMOXOR.W
01100	aq	rl	rs2	rs1	010	rd	0101111	AMOAND.W
01000	aq	rl	rs2	rs1	010	rd	0101111	AMOOR.W
10000	aq	rl	rs2	rs1	010	rd	0101111	AMOMIN.W
10100	aq	rl	rs2	rs1	010	rd	0101111	AMOMAX.W
11000	aq	rl	rs2	rs1	010	rd	0101111	AMOMINU.W
11100	aq	rl	rs2	rs1	010	rd	0101111	AMOMAXU.W

RV64A Standard Extension (in addition to RV32A)								
00010	aq	rl	00000	rs1	011	rd	0101111	LR.D
00011	aq	rl	rs2	rs1	011	rd	0101111	SC.D
00001	aq	rl	rs2	rs1	011	rd	0101111	AMOSWAP.D
00000	aq	rl	rs2	rs1	011	rd	0101111	AMOADD.D
00100	aq	rl	rs2	rs1	011	rd	0101111	AMOXOR.D
01100	aq	rl	rs2	rs1	011	rd	0101111	AMOAND.D
01000	aq	rl	rs2	rs1	011	rd	0101111	AMOOR.D
10000	aq	rl	rs2	rs1	011	rd	0101111	AMOMIN.D
10100	aq	rl	rs2	rs1	011	rd	0101111	AMOMAX.D
11000	aq	rl	rs2	rs1	011	rd	0101111	AMOMINU.D
11100	aq	rl	rs2	rs1	011	rd	0101111	AMOMAXU.D

- LR/SC: “Load Reserved/Store Conditional” on a memory location
- AMOxxx: atomic read-modify-writes on a memory location
- W ops operate on aligned 32-bit words
- D ops operate on aligned 64-bit words

AMO instructions (Atomic Memory Ops)



- LR/SC: Typically in a loop, until success
- LR/SC: Many nuances on “address region”, “observed by another hart”, number and type of instructions between the LR and SC, ...

* “hart” = “hardware thread”, a single hardware thread, based on a single instruction-fetch unit

“F” extension: IEEE Single-precision Floating Point

State: 32 x 32-bit floating-point registers, plus these CSRs (Control/Status Regs) holding standard IEEE flags:

Number	Privilege	Name	Description
Floating-Point Control and Status Registers			
0x001	Read/write	fflags	Floating-Point Accrued Exceptions.
0x002	Read/write	frm	Floating-Point Dynamic Rounding Mode.
0x003	Read/write	fcsr	Floating-Point Control and Status Register (frm + fflags).

Instructions:

RV32F Standard Extension							
imm[11:0]			rs1	010	rd	0000111	FLW
imm[11:5]		rs2	rs1	010	imm[4:0]	0100111	FSW
rs3	00	rs2	rs1	rm	rd	1000011	FMADD.S
rs3	00	rs2	rs1	rm	rd	1000111	FMSUB.S
rs3	00	rs2	rs1	rm	rd	1001011	FNMSUB.S
rs3	00	rs2	rs1	rm	rd	1001111	FNMADD.S
0000000		rs2	rs1	rm	rd	1010011	FADD.S
0000100		rs2	rs1	rm	rd	1010011	FSUB.S
0001000		rs2	rs1	rm	rd	1010011	FMUL.S
0001100		rs2	rs1	rm	rd	1010011	FDIV.S
0101100		00000	rs1	rm	rd	1010011	FSQRT.S
0010000		rs2	rs1	rm	rd	1010011	FSCN1.S
(... more ... please consult spec)							

RV64F Standard Extension (in addition to RV32F)						
1100000	00010	rs1	rm	rd	1010011	FCVT.L.S
1100000	00011	rs1	rm	rd	1010011	FCVT.LU.S
1101000	00010	rs1	rm	rd	1010011	FCVT.S.L
1101000	00011	rs1	rm	rd	1010011	FCVT.S.LU

“D” extension: IEEE Double-precision Floating Point

State: 32 x 64-bit floating-point registers, plus these CSRs (Control/Status Regs) holding standard IEEE flags:

Number	Privilege	Name	Description
Floating-Point Control and Status Registers			
0x001	Read/write	fflags	Floating-Point Accrued Exceptions.
0x002	Read/write	frm	Floating-Point Dynamic Rounding Mode.
0x003	Read/write	fcsr	Floating-Point Control and Status Register (frm + fflags).

Instructions:

RV32D Standard Extension							
imm[11:0]			rs1	011	rd	0000111	FLD
imm[11:5]		rs2	rs1	011	imm[4:0]	0100111	FSD
rs3	01	rs2	rs1	rm	rd	1000011	FMADD.D
rs3	01	rs2	rs1	rm	rd	1000111	FMSUB.D
rs3	01	rs2	rs1	rm	rd	1001011	FNMSUB.D
rs3	01	rs2	rs1	rm	rd	1001111	FNMADD.D
0000001		rs2	rs1	rm	rd	1010011	FADD.D
0000101		rs2	rs1	rm	rd	1010011	FSUB.D
0001001		rs2	rs1	rm	rd	1010011	FMUL.D
0001101		rs2	rs1	rm	rd	1010011	FDIV.D
0101101		00000	rs1	rm	rd	1010011	FSQRT.D
0010001		rs2	rs1	rm	rd	1010011	FSCN.D
(... more ... please consult spec)							

RV64D Standard Extension (in addition to RV32D)						
1100001	00010	rs1	rm	rd	1010011	FCVT.L.D
1100001	00011	rs1	rm	rd	1010011	FCVT.LU.D
1110001	00000	rs1	000	rd	1010011	FMV.X.D
1101001	00010	rs1	rm	rd	1010011	FCVT.D.L
1101001	00011	rs1	rm	rd	1010011	FCVT.D.LU
1111001	00000	rs1	000	rd	1010011	FMV.D.X

“C” extension: Compressed Instructions for smaller footprint

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	
000		0										0		00		<i>Illegal instruction</i>
000		nzuimm[5:4 9:6 2 3]										rd'		00		C.ADDI4SPN (<i>RES</i> , nzuimm=0)
001		uimm[5:3]			rs1'			uimm[7:6]			rd'		00		C.FLD (<i>RV32/64</i>)	
001		uimm[5:4 8]			rs1'			uimm[7:6]			rd'		00		C.LQ (<i>RV128</i>)	
010		uimm[5:3]			rs1'			uimm[2 6]			rd'		00		C.LW	
011		uimm[5:3]			rs1'			uimm[2 6]			rd'		00		C.FLW (<i>RV32</i>)	
011		uimm[5:3]			rs1'			uimm[7:6]			rd'		00		C.LD (<i>RV64/128</i>)	
100		—												00		<i>Reserved</i>
101		uimm[5:3]			rs1'			uimm[7:6]			rs2'		00		C.FSD (<i>RV32/64</i>)	
101		uimm[5:4 8]			rs1'			uimm[7:6]			rs2'		00		C.SQ (<i>RV128</i>)	
110		uimm[5:3]			rs1'			uimm[2 6]			rs2'		00		C.SW	
111		uimm[5:3]			rs1'			uimm[2 6]			rs2'		00		C.FSW (<i>RV32</i>)	
111		uimm[5:3]			rs1'			uimm[7:6]			rs2'		00		C.SD (<i>RV64/128</i>)	

Table 16.5: Instruction listing for RVC, Quadrant 0.

(... more ... please consult spec)

- “C” instructions are 16-bits wide, can be packed two-to-a-32b-word.
- Not standalone; are mixed with standard 32-bit instructions.
- 3-bit register fields refer to the 8 “most popular” registers.
- Each 16-bit C instruction expands to a standard 32-bit RV32/64 instruction; so only needs hardware at the front-end of the CPU pipeline (16-bit fetch and expansion).
- Include M, F, D instructions.

Notes on Memory and I/O

- Pure “Load-Store” ISA (with register base-address + index), *i.e.*, total separation of memory instructions vs. non-memory instructions (computation, system, ...).
- Flat memory space (address-width depends on RV32 or RV64, and Virtual Memory Scheme).
- All I/O through memory-mapped device registers; no separate I/O instructions.
- Optional PMP CSRs (Physical Memory Protection) can impose base-and-bounds regions on memory with access permissions. This provides lightweight sandboxing without the overheads of full Virtual Memory (TLBs/MMUs, page tables, address translation, ...).
- FENCE, FENCE.I, and SFENCE.VMA instructions for implementations that need to synchronize multicores, I-Caches and D-Caches, and TLBs/MMUs with Caches.

“Weak Memory Models” for concurrent access (in multithreading, multicore)

In multithreaded implementations, when multiple harts perform concurrent memory operations (fetches, loads, stores, AMOOps), the order in which results are returned to each hart can depend on micro-architectural performance features such as reordering, caches, cache-line size, store-buffers, *etc.*. Without a formal weak memory model, this can be chaotic and it becomes impossible/very hard to write reliable multithreaded software.

An ISA’s *weak memory model* defines what are legal outcomes for concurrent accesses, across all implementations. Every major ISA also has a defined weak memory model.

In the RISC-V ISA:

- RISC-V Weak Memory Model (similar to ARMv8).
- Optional TSO (Total Store Order) memory model (like x86, Sparc).

TSO is a stronger memory model, and much legacy code is written assuming this model (originating on x86, Sparc).

Code written for TSO can run on a RISC-V implementation that only implements the Weak Memory Model, by inserting suitable “fence” instructions and ordering directives into the code. In many cases, this can involve just recompiling libraries for concurrent data structures.

Standard Privileged ISA

* Because of clean ECALL from Unprivileged ISA, easy to substitute non-standard Privileged ISAs, *e.g.*, containing just basic trap handling for embedded applications.

Standard Privileged ISA

Note: The separation between Unprivileged and Privileged ISAs is very clean. A CPU implementation can easily implement a non-standard Privileged ISA with the standard Unprivileged ISA, if desired.

Privilege Levels: enables clean *virtualization* (hypervisors, Xen, VMWare, ...)

Level	Code	Name	Comment
3	11	M (Machine)	Highest privilege; firmware, boot loaders, ... Operating systems (like Linux) Applications
2	10	<i>Reserved</i>	
1	01	S (Supervisor)	
0	00	U (User/Application)	


Implementations do not have to implement all privilege levels (implementation cost tradeoff).
Standard combinations:

Implemented Levels	Intended Usage
M	Simple embedded systems (“bare metal”) Secure embedded systems Full OS + apps
M,U	
M,S,U	

Standard Privileged ISA: CSRs (Control and Status Registers)

- CSRs: additional set of registers in the CPU, identified by 12-bit CSR number. Thus, up to 4096 possible CSRs, but most are optional; only a small handful are essential in an implementation.
- CSRs can be accessed programmatically with the CSR instructions shown below, each of which atomically swaps the contents of a CSR with standard integer registers, with some variations and nuances.

31	20 19	15 14	12 11	7 6	0
csr	rs1	funct3	rd	opcode	
12	5	3	5	7	
source/dest	source	CSRRW	dest	SYSTEM	
source/dest	source	CSRRS	dest	SYSTEM	
source/dest	source	CSRRC	dest	SYSTEM	
source/dest	uimm[4:0]	CSRRWI	dest	SYSTEM	
source/dest	uimm[4:0]	CSRRSI	dest	SYSTEM	
source/dest	uimm[4:0]	CSRRCI	dest	SYSTEM	



- The upper 4 bits of a CSR's number specifies which privilege levels can access that particular CSR, and with which ops (read, read/write).
- Some CSRs are “shadows” of other CSRs: when viewed at lower privilege levels, only some bits may be visible, or may be read-only vs. read-write, *etc.*.

Example: `cycle` (User privilege) is a read-only shadow of `mcycle` (Machine privilege).

CSRs accessible from User level

Number	Privilege	Name	Description
Unprivileged Floating-Point CSRs			
0x001	URW	<code>fflags</code>	Floating-Point Accrued Exceptions.
0x002	URW	<code>frm</code>	Floating-Point Dynamic Rounding Mode.
0x003	URW	<code>fcsr</code>	Floating-Point Control and Status Register (<code>frm</code> + <code>fflags</code>).
Unprivileged Counter/Timers			
0xC00	URO	<code>cycle</code>	Cycle counter for RDCYCLE instruction.
0xC01	URO	<code>time</code>	Timer for RDTIME instruction.
0xC02	URO	<code>instret</code>	Instructions-retired counter for RDINSTRET instruction.
0xC03	URO	<code>hpmcounter3</code>	Performance-monitoring counter.
0xC04	URO	<code>hpmcounter4</code>	Performance-monitoring counter.
		⋮	
0xC1F	URO	<code>hpmcounter31</code>	Performance-monitoring counter.
0xC80	URO	<code>cycleh</code>	Upper 32 bits of <code>cycle</code> , RV32 only.
0xC81	URO	<code>timeh</code>	Upper 32 bits of <code>time</code> , RV32 only.
0xC82	URO	<code>instreth</code>	Upper 32 bits of <code>instret</code> , RV32 only.
0xC83	URO	<code>hpmcounter3h</code>	Upper 32 bits of <code>hpmcounter3</code> , RV32 only.
0xC84	URO	<code>hpmcounter4h</code>	Upper 32 bits of <code>hpmcounter4</code> , RV32 only.
		⋮	
0xC9F	URO	<code>hpmcounter31h</code>	Upper 32 bits of <code>hpmcounter31</code> , RV32 only.

Table 2.2: Currently allocated RISC-V unprivileged CSR addresses.

- CSRs for floating point
- CSRs for timing: real-time, cycle, instruction count
- CSRs for other performance counters (implementation-defined)

CSRs accessible from Supervisor level

Number	Privilege	Name	Description
Supervisor Trap Setup			
0x100	SRW	sstatus	Supervisor status register.
0x104	SRW	sie	Supervisor interrupt-enable register.
0x105	SRW	stvec	Supervisor trap handler base address.
0x106	SRW	scounteren	Supervisor counter enable.
Supervisor Configuration			
0x10A	SRW	senvcfg	Supervisor environment configuration register.
Supervisor Trap Handling			
0x140	SRW	sscratch	Scratch register for supervisor trap handlers.
0x141	SRW	sepc	Supervisor exception program counter.
0x142	SRW	scause	Supervisor trap cause.
0x143	SRW	stval	Supervisor bad address or instruction.
0x144	SRW	sip	Supervisor interrupt pending.
Supervisor Protection and Translation			
0x180	SRW	satp	Supervisor address translation and protection.
Debug/Trace Registers			
0x5A8	SRW	scontext	Supervisor-mode context register.

Table 2.3: Currently allocated RISC-V supervisor-level CSR addresses.

- CSRs for handling exceptions (traps and interrupts) at Supervisor level (see later slides)
- **satp**: for Virtual memory (see later slides)


CSRs accessible from Machine level

Number	Privilege	Name	Description
Machine Information Registers			
0xF11	MRO	<code>mvendorid</code>	Vendor ID.
0xF12	MRO	<code>marchid</code>	Architecture ID.
0xF13	MRO	<code>mimpid</code>	Implementation ID.
0xF14	MRO	<code>mhartid</code>	Hardware thread ID.
0xF15	MRO	<code>mconfigptr</code>	Pointer to configuration data structure.
Machine Trap Setup			
0x300	MRW	<code>mstatus</code>	Machine status register.
0x301	MRW	<code>misa</code>	ISA and extensions
0x302	MRW	<code>medeleg</code>	Machine exception delegation register.
0x303	MRW	<code>mideleg</code>	Machine interrupt delegation register.
0x304	MRW	<code>mie</code>	Machine interrupt-enable register.
0x305	MRW	<code>mtvec</code>	Machine trap-handler base address.
0x306	MRW	<code>mcounteren</code>	Machine counter enable.
0x310	MRW	<code>mstatush</code>	Additional machine status register, RV32 only.
Machine Trap Handling			
0x340	MRW	<code>mscratch</code>	Scratch register for machine trap handlers.
0x341	MRW	<code>mepc</code>	Machine exception program counter.
0x342	MRW	<code>mcause</code>	Machine trap cause.
0x343	MRW	<code>mtval</code>	Machine bad address or instruction.
0x344	MRW	<code>mip</code>	Machine interrupt pending.
0x34A	MRW	<code>mtinst</code>	Machine trap instruction (transformed).
0x34B	MRW	<code>mtval2</code>	Machine bad guest physical address.

(... more ... please consult spec)

- CSRs for configuration discovery from software
- CSRs for handling exceptions (traps and interrupts) at Machine level (see later slides)

Exceptions (Traps and Interrupts)



- **mtvec** (exception handler address) has been pre-loaded before this scenario (by boot-loader, OS, ...)
- **mepc** may be updated by the trap handler to $PC3 = PC1+0$, $PC1+2$, $PC1+4$, or something else entirely, depending on whether the trapping instruction needs to be retried or skipped, whether it is a compressed instruction (+2) or not (+4), and whether we're resuming the trapping code or switching to an entirely different thread (*e.g.*, OS thread-switching after time-slice).
- “Arguments” and “results” are passed in registers, according to a calling convention/ABI (Application Binary Interface).
- Can be nested, *i.e.*, exception handler may itself encounter trap/interrupt.
- Exceptions can be delivered at Machine and Supervisor privilege levels.
- The ISA defines mechanisms to automatically *delegate* exception-handling to a lower privilege level (Supervisor, User); see **medeleg** and **mideleg** CSRs.

CSR MStatus




Figure 3.7: Machine-mode status register (`mstatus`) for RV64.

- `mstatus` is quite central and critical, and manipulation of its bits involves some complexity and subtlety; we recommend reading the spec very closely and carefully.
- Various bits represent a 2-level “stack” (indexed by privilege level) that are conceptually pushed on exception entry and popped on returns from exceptions (see “update status” on previous slide). These include interrupt enable bits (MIE/MPIE, SIE/SPIE), previous privilege level (MPP, SPP).
- The RV32 version of `mstatus` omits some of the bits.
- There is also an `sstatus` CSR at Supervisor level, which is similar to, but only a subset of `mstatus`.

CSR MCause

Interrupt	Exception Code	Description
1	0	<i>Reserved</i>
1	1	Supervisor software interrupt
1	2	<i>Reserved</i>
1	3	Machine software interrupt
1	4	<i>Reserved</i>
1	5	Supervisor timer interrupt
1	6	<i>Reserved</i>
1	7	Machine timer interrupt
1	8	<i>Reserved</i>
1	9	Supervisor external interrupt
1	10	<i>Reserved</i>
1	11	Machine external interrupt
1	12–15	<i>Reserved</i>
1	≥16	<i>Designated for platform use</i>
0	0	Instruction address misaligned
0	1	Instruction access fault
0	2	Illegal instruction
0	3	Breakpoint
0	4	Load address misaligned
0	5	Load access fault
0	6	Store/AMO address misaligned
0	7	Store/AMO access fault
0	8	Environment call from U-mode
0	9	Environment call from S-mode
0	10	<i>Reserved</i>
0	11	Environment call from M-mode
0	12	Instruction page fault
0	13	Load page fault
0	14	<i>Reserved</i>
0	15	Store/AMO page fault
0	16–23	<i>Reserved</i>
0	24–31	<i>Designated for custom use</i>
0	32–47	<i>Reserved</i>
0	48–63	<i>Designated for custom use</i>
0	≥64	<i>Reserved</i>

- On any exception, the cause is recorded in CSR `mcause` during the transfer of control to the exception handler.
- Exception-handler code uses `mcause` to discover what kind of event caused this exception, using which it can branch to event-type-specific handlers.

CSRs MIP and MIE

15	12	11	10	9	8	7	6	5	4	3	2	1	0
0	MEIP	0	SEIP	0	MTIP	0	STIP	0	MSIP	0	SSIP	0	0
4	1	1	1	1	1	1	1	1	1	1	1	1	1

Figure 3.14: Standard portion (bits 15:0) of **mip**.

15	12	11	10	9	8	7	6	5	4	3	2	1	0
0	MEIE	0	SEIE	0	MTIE	0	STIE	0	MSIE	0	SSIE	0	0
4	1	1	1	1	1	1	1	1	1	1	1	1	1

Figure 3.15: Standard portion (bits 15:0) of **mie**.

- Interrupts arrive at a RISC-V hart only via the **mip** CSR.
- Both are full-width registers (32b in RV32, 64b in RV64) but only the bottom 16 bits have standard definitions.
- **mip**: Machine-level register for Interrupts-pending (there is also an **sip** for Supervisor level). The bits can be set by various sources of interrupts outside the CPU (see other slides on CLIC and PLIC for examples).
- Sources can be external devices (“E”), timers (“T”) and inter-processor software interrupts (“S”).
- **mie**: Machine-level register for Interrupts-enabled (there is also an **sie** for Supervisor level). The CPU can mask-out specific source of interrupts by writing 0 to the corresponding bit.

Virtual Memory


- Virtual Memory (VM) is available in Supervisor and User privilege levels. Most Linux OS code runs at Supervisor level.
- VM is optional; an implementation can have Supervisor and User levels without implementing VM.
- The Privileged ISA spec defines several possible virtual memory schemes:
 - For RV32: *Sv32*
 - For RV64: *Sv39*, *Sv48* and *Sv57*

(XX in SvXX indicates the width of virtual addresses.)

Next two slides illustrate Sv32 and Sv39, respectively, the two most common variants.


- Implementations can choose *how* to implement VM, in particular VA→PA translation:
 - in hardware or firmware or trap-handlers;
 - whether to use accelerators like TLBs (Translation Look-aside Buffers),
 - *etc.*
 - The SFENCE.VMA instruction is available to synchronize updates to address-translation hardware.

Virtual Memory Page Table for RV32/Sv32 using CSR `satp`



- When running in VM mode (indicated by `mstatus` bits, privilege level, *etc.*), the PC, and addresses in LD/ST instructions are VAs (Virtual Addresses).
- CSR `satp` contains a pointer to a 2-level tree. Each node in the tree is a *PTN* (Page Table Node), an aligned 4KiB block containing 1024 *PTEs* (Page Table Entries, each 32b). Bits in each PTE indicate whether it is valid/invalid, a leaf (pointing to a data page), or a pointer to the next-level PTN.
- Leaves at level 1 point at 2MiB naturally aligned “megapages”; leaves at level 0 point at 4KiB naturally aligned “pages”;
- Sv32 addresses have 32 bits. 10-bit fields are used to index a PTE in a PTN, and the 12 LSBs are used to index a byte in a 4KiB page.
- If a VA→PA translation fails (encounter invalid PTE, protection failure, ...) it raises a *page fault* exception or an *access fault* exception, which is handled in the usual way (see “Exceptions” slide earlier).

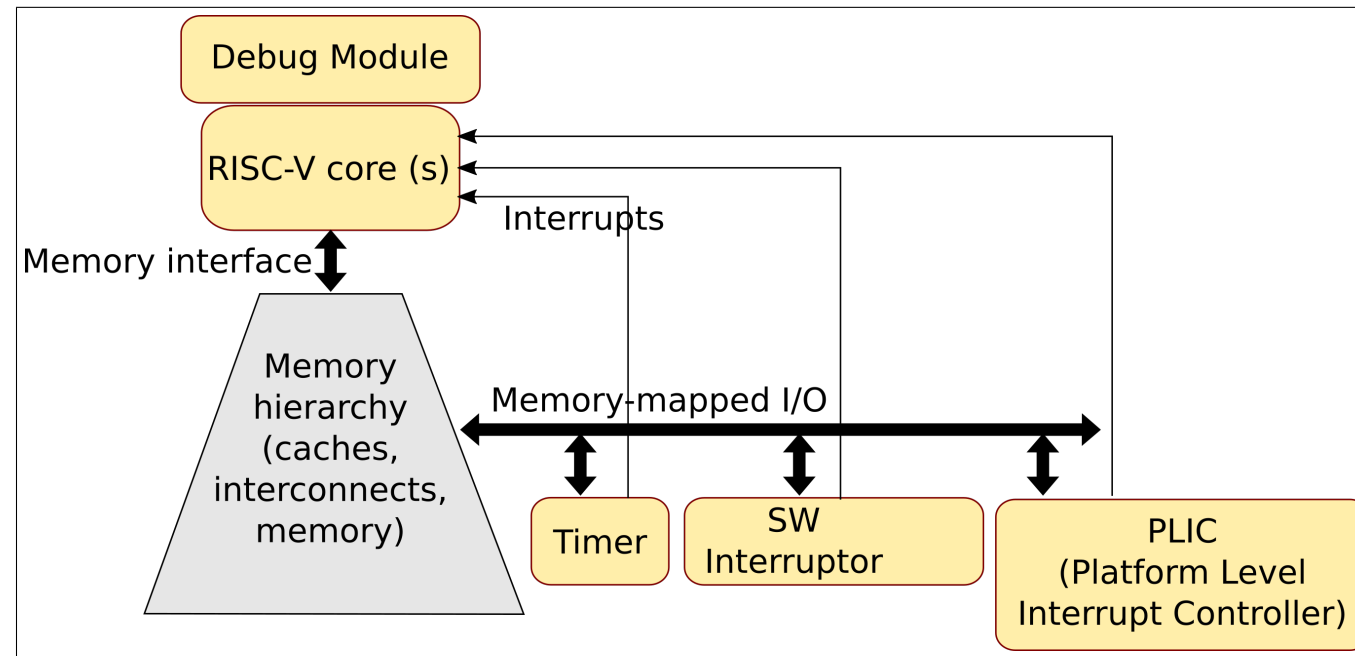
Virtual Memory Page Table for RV64/Sv39 using CSR `satp`



- When running in VM mode (indicated by `mstatus` bits, privilege level, *etc.*), the PC, and addresses in LD/ST instructions are VAs (Virtual Addresses).
- CSR `satp` contains a pointer to a 3-level tree. Each node in the tree is a *PTN* (Page Table Node), an aligned 4KiB block containing 512 *PTEs* (Page Table Entries, each 64b). Bits in each PTE indicate whether it is valid/invalid, a leaf (pointing to a data page), or a pointer to the next-level PTN.
- Leaves at level 2 point at 1GiB naturally aligned “gigapages”; leaves at level 1 point at 2MiB naturally aligned “megapages”; leaves at level 0 point at 4KiB naturally aligned “pages”;
- Sv39 addresses have 39 bits (bits above this are unused). 9-bit fields are used to index a PTE in a PTN, and the 12 LSBs are used to index a byte in a 4KiB page.
- If a VA→PA translation fails (encounter invalid PTE, protection failure, ...) it raises a *page fault* exception or an *access fault* exception, which is handled in the usual way (see “Exceptions” slide earlier).

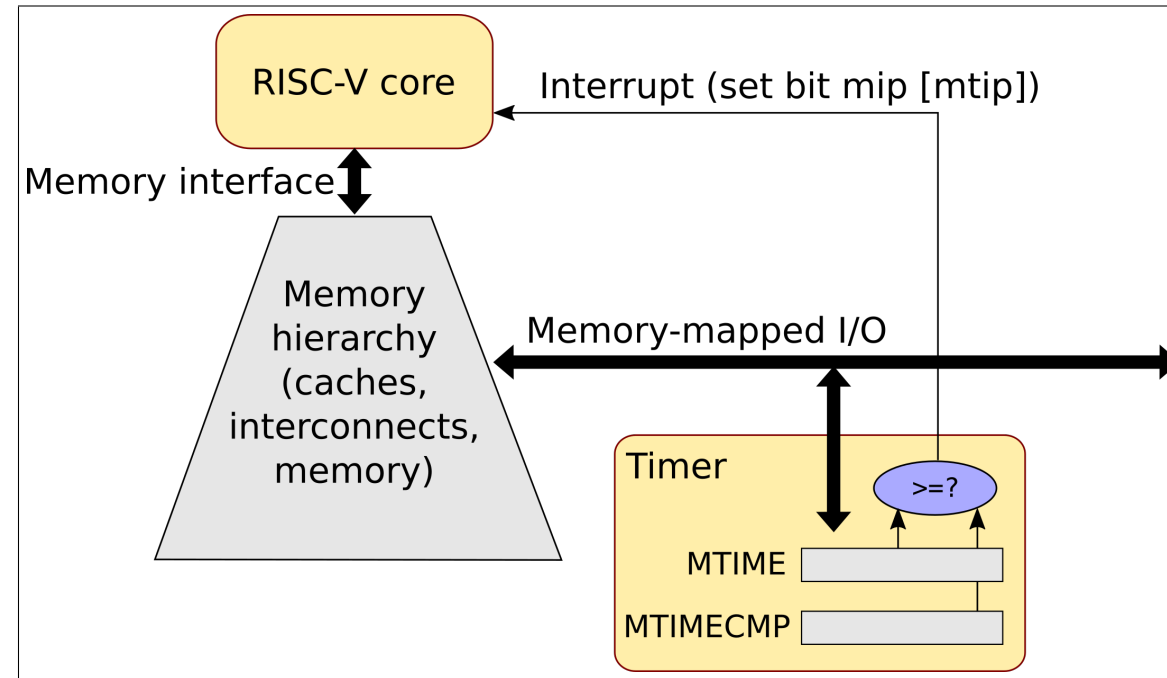
Common System Components adjunct to RISC-V cores

Common System Components adjunct to RISC-V Cores




Each is described in more detail in the following slides.

Timer: Common System Component accompanying RISC-V Cores



- For SW to measure real-time, and for real-time timer interrupts.
- Contains two memory-mapped registers, MTIME and MTIMECMP.
- MTIME ticks upwards constantly (can also be written by CPU to initialize).
- Timer Interrupts: typical usage:
 - CPU reads MTIME (let's call it t)
 - CPU writes $t + \delta$ in MTIMECMP
 - When MTIME ticks up by δ (so $\text{MTIME} \geq \text{MTIMECMP}$), delivers an interrupt to CSR MIP at the MTIP bit position.

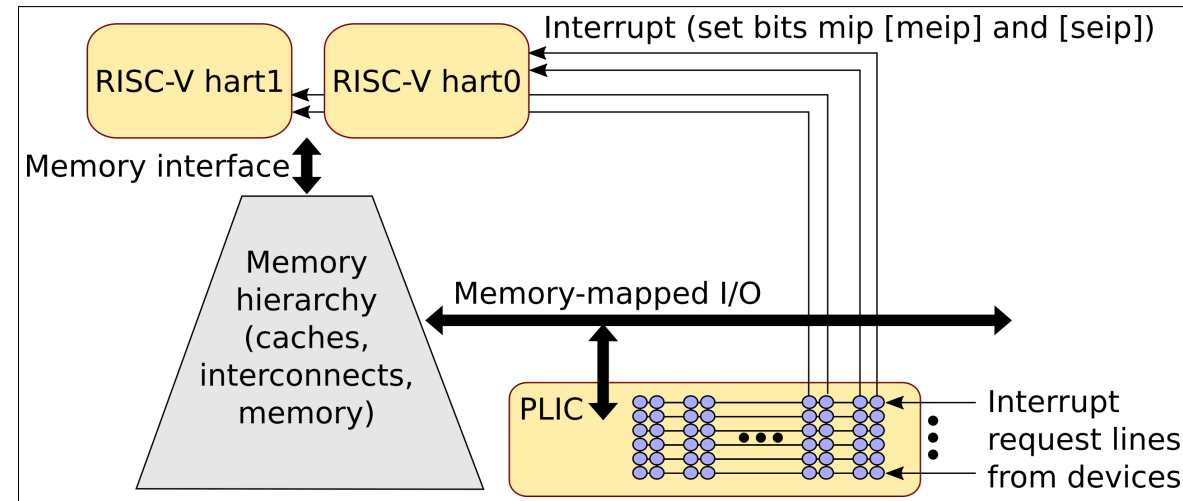
SW Interruptor: Common System Component accompanying RISC-V Cores



- For a hart to deliver a “software interrupt” to another (or same) hart.
- Illustration shows only two harts, but there could be more (multicore).
- Contains one memory-mapped register, MSIP
- SW Interruptor: typical usage:
 - Hart I writes MSIP [J] to deliver a SW interrupt to hart J
 - MSIP [J] is connected to the CSR MIP at the MSIP bit position (in core J).

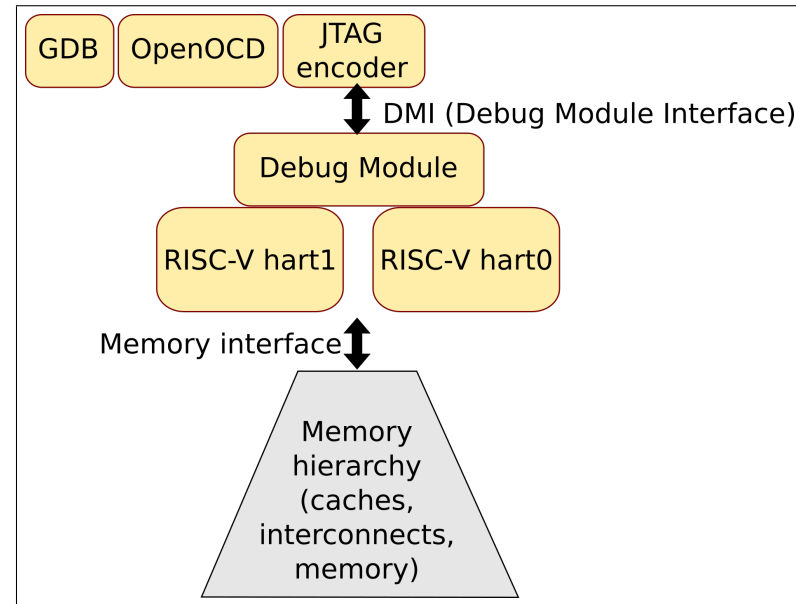
In some systems MTIME, MTIMECMP, MSIP are in a common component called the *CLINT* (Core-Local Interruptor).

PLIC: Common System Component accompanying RISC-V Cores



- Conceptually a matrix connecting an interrupt request line from a device to MIP[MEIP] or MIP[SEIP] of any of multiple harts.
- Illustration shows only two harts, but there could be more (multicore).
- Matrix nodes can be programmed by MMIO from cores to configure connectivity, priority, arbitration, interrupt masks, *etc.*
- A device interrupt can connect to multiple cores; PLIC has facilities for competing core to “claim” an interrupt atomically.

Debug Module: Common System Component accompanying RISC-V Cores



- A RISC-V Debug Module is a HW module closely connected to multiple harts. Can control them individually and collectively.
- Functions: Reset hart, read/write registers and CSRs, read/write memory, stop/start/continue/single-step, set breakpoints, set watchpoints, ...
- Is controlled through a standard memory-like Debug Module Interface (DMI).
- DMI can be connected to JTAG (open-source RTL available).
- JTAG can connect to standard, open-source OpenOCD (Open On-Chip Debugger).
- DMI or OpenOCD can connect to standard GDB.

With the Debug Module, a hardware RISC-V core can be controlled precisely from GDB and/or OpenOCD, starting at the very first instruction out of reset.

A Brief Note about the RISC-V Software Ecosystem

The following SW development tools have been available for many years and are mature (upstreamed into standard distributions):

- GNU toolchain (gcc, gdb, ...)
- Linux, FreeRTOS, seL4, ...

Hundreds/thousands of other porting efforts are underway (operating systems, programming language implementations, applications, tools), in hundreds/thousands of companies and universities worldwide; the ecosystem grows every day.

HW Verification of CPU and System Implementations

- Formal and Reference Models of RISC-V CPUs
- Standard ISA Tests
- Google open source “riscv-dv” instruction stream generator

Formal Specification of RISC-V ISA

RISC-V is the *only* major ISA with a public formal specification.*

Formal Specification (Unprivileged RV32/64 IMAFDC, Privileged):

- The formal spec itself: <https://github.com/riscv/sail-riscv>
- “A Tour of the RISC-V ISA Formal Specification”: https://github.com/rsnikhil/RISCV_ISA_Spec_Tour by R.S.Nikhil, RISC-V Summit 2019. Video: <https://www.youtube.com/watch?v=k3NhEtk8TAs>

Written in **Sail**, a language (DSL) for ISA formal specifications from U.Cambridge, UK:

- Sail has also been used for a full formal spec for ARMv8 ISA, and for partial formal specs for Power, MIPs and x86.
- Sail website: <https://github.com/rem-s-project/sail>.
- Tutorial: “High Level Sail Overview”, by Bill McSpadden, RISC-V Summit 2022.

The formal spec can be compiled into a simulator for RISC-V, which can be used as a golden reference model.

* We believe ARMv8 also has a formal specification, but only available internally inside ARM.

Golden Reference Models of the RISC-V ISA

Software simulators:

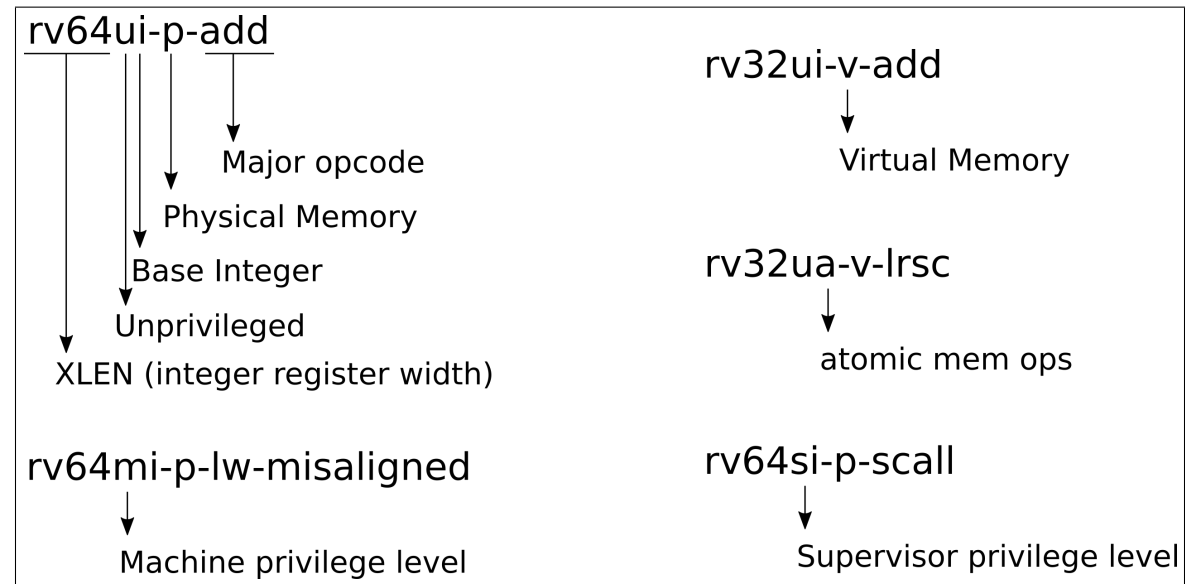
- These are pure simulators of the ISA, and do not simulate any micro-architectural details (pipelining, caches, MMUs, ...). They can produce an instruction trace that can be compared with the instruction trace from a hardware implementation.
- SPIKE (The “Golden Spike”) is the principal simulator reference model written in C++
<https://github.com/riscv-software-src/riscv-isa-sim>
- Other:
 - Sail RISC-V formal spec, compiled to an executable.
 - In QEMU, TinyEMU, Imperas, ... (*many*)
 - Open-source formal spec written in Haskell: https://github.com/rsnikhil/Forvis_RISCV-ISA-Spec

Hardware reference models:

- These are synthesizable HW implementations of RISC-V written for clarity (correctness), not performance, and can be run in HW (typically in FPGAs or HW simulation), adjacent to a real HW implementation, to compare ISA traces directly (“Tandem Verification”).
- Examples:
 - RVBS, from U.Cambridge, UK: <https://github.com/CTSRD-CHERI/RVBS>
 - Magritte, from Bluespec, Inc.

Standard ISA Tests

- Reference: <https://github.com/riscv-software-src/riscv-tests>
- These are a set of several hundred programs, written in RISC-V Assembly language.
- Each program focuses on one particular opcode, and contains several test 1..N.
 - Each test ends with a postlude that either signals success, or failure with an indication of which of its N tests first failed.
 - The result signal is to store a value to a memory-mapped `tohost` location. Different implementations implement this in different ways.
- The test programs are named according to the ISA modular structure. Examples:



Other Verification Tests

Google open source “riscv-dv” instruction stream generator

- Reference: <https://github.com/google/riscv-dv>
“RISCV-DV is a SV/UVM based open-source instruction generator for RISC-V processor verification.”
- Typically run on reference golden model, and on DUT, and compare instruction traces.

Beyond ISA tests, HW verification is often performed by comparing behavior of reference golden model against HW implementation while booting an OS, typically Linux.

Would be useful: Test code whose execution time in simulation is between ISA tests (seconds) and Linux kernel (days), gradually adding virtual memory, interrupts, and process-switching:

- ISA tests’ execution time: seconds (even in simulation).
- Linux kernel execution time: Many hours/days in simulation.

Perhaps a microkernel (seL4?) may serve this purpose.

End

Please direct questions/comments/errata to the author.

Possible topics to add in future

- Verification under non-determinism (*e.g.*, with timer interrupts; multicore)
- Tandem Verification
- Verification of Weak Memory Model implementations
- Formal Verification